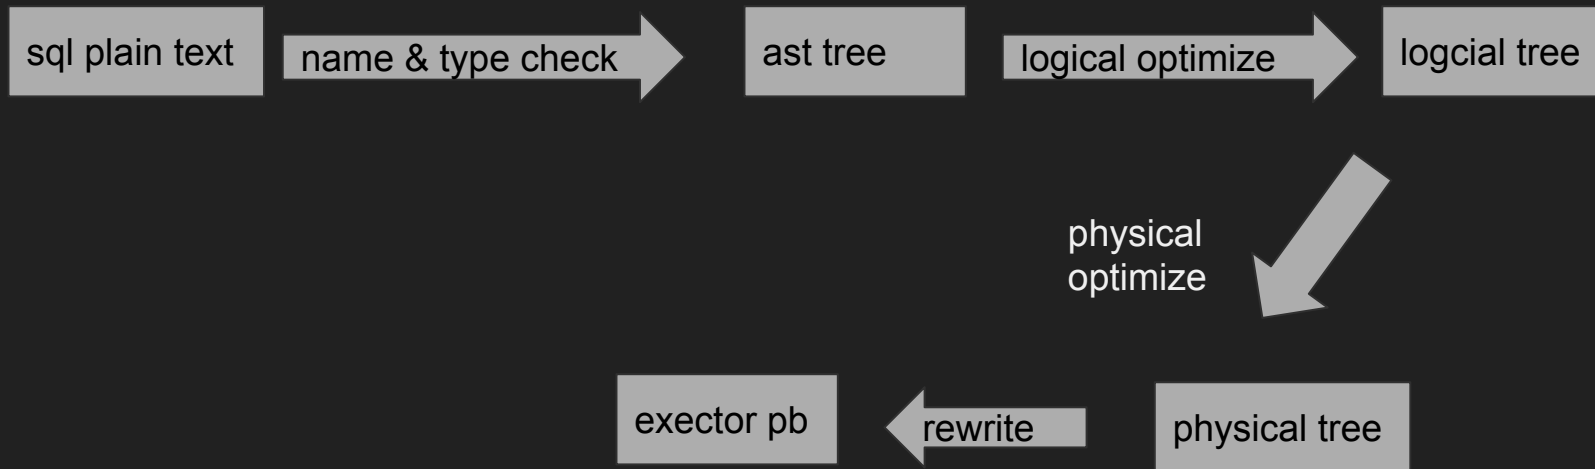


An Overview of Query Optimization

Han Fei@PingCAP

A optimizer paradigm

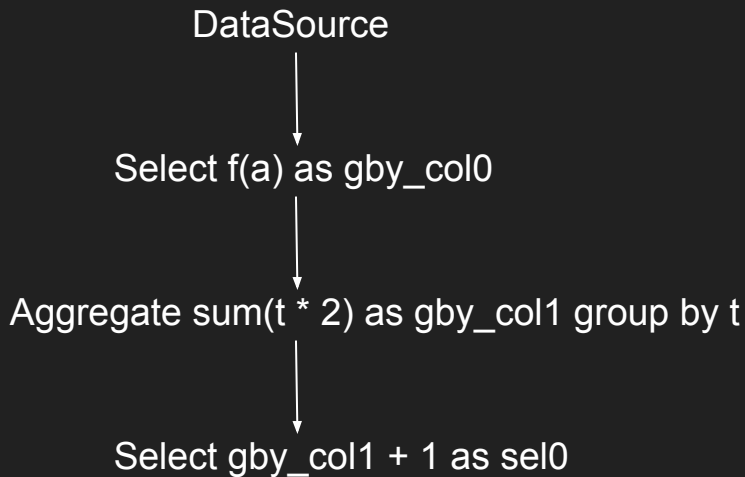


logical optimize

- common expression elimination
- column pruning
- constant fold
- constant propagation
- expression minimize
- predicate/projection/aggregation pushdown

Common Expression Elimination

- Example: Select $\text{sum}(f(a) * 2) + 1$ from src group by $f(a)$
- Can be converted to:



How to identify common expression ?

Calculate Hash Code

Predicate Pushdown

- convert DNF to CNF
 - Example: `A JOIN B WHERE A.id > 10 or (A.id < 5 and b.id < 5)`
- split CNF items
- when predicate meets projection, aggregation, join, union, etc.
- the opt can do during ppd:
 - outer join elimination
 - Example: `A LEFT OUTER JOIN B WHERE B.id < 5`
 - constant propagation
 - Example: `A JOIN B ON A.id = B.id WHERE A.id < 5`
 - Example: `A JOIN B ON A.id = B.id JOIN C ON B.id = C.id WHERE A.id < 5`

physical algebra

- stream aggr & hash aggr
- index join, merge join, hash join
- subqueries, semi join
 - `select * from t where t.id in (select id from s)`
 - `select * from t where t.id in (select count(*) from s group by s.key)`
 - `select * from t where t.id in (select count(*) from s group by s.key having s.id = t.key)`
 - `select * from t where t.id in (select count(*) from s group by s.key where s.id = t.key)`
 - `select count(*) from t group by t.id in (select count(*) from s group by s.key having s.id = t.key)`

aggregation/distinct push down

- aggregate meets join
 - `SELECT SUM(DISTINCT t.id) FROM t JOIN s ON t.id = s.id`
 - `SELECT SUM(DISTINCT t.id) FROM t JOIN s ON t.key = s.key`
- aggregate meets union all
 - `SELECT SUM(tmp.id) FROM (select * FROM T) UNION ALL (SELECT * FROM S)`
 - `SELECT AVG(tmp.id) FROM (select * FROM T) UNION ALL (SELECT * FROM S)`

Physical Optimize

- Cost Based Optimization
- History Based Optimization
- Data Skew Optimization

Cost Modal

- Combines components of estimated
 - CPU (instructions)
 - I/O (random and sequential)
 - Communications between nodes
- Basic Statistics
 - Number of rows in tables
 - for each columns
 - distinct value, avg length of data values, data range infos (histogram)

Interesting Order, Group

- order
 - `SELECT * FROM A JOIN B ON A.ID = B.ID ORDER BY A.KEY`
- group
 - `SELECT SUM(A.KEY) FROM A JOIN B ON A.ID = B.ID GROUP BY A.ID`
 - A , B Sort by ID and apply merge join
 - A Group by ID and apply hash join

Join ReOrder

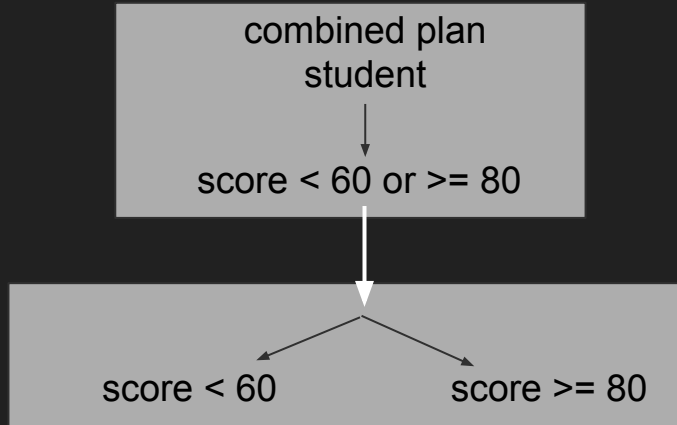
- Specify the EquivalenceClass
 - a join b on a.id = b.id join c on a.key = c.key join d on a.id = d.id
 - a b d can use merge join in one task
- For a few join items : Dynamic Programming
- Else apply a greedy algo

Advised Topic: History based optimization

- statistics feedback and self-adapt

Advised Topic: Common queries(tables) combine

- `SELECT id from student where score < 60`
- `SELECT id from student where score >= 80`
- Two Choices:



Skew Aggregation, Join

Search Engine

- Bottom-Up (R system -> Starburst, DB2, Oracle)
- TopDown (Volcano -> Cascades, SQL-Server, SCOPE, Calcite)

Thanks!